

# Your Title Here

**Name(s):** (your name(s) here)

**Website Link:** (your website link)

```
In [ ]: import pandas as pd
import numpy as np
from pathlib import Path

import plotly.express as px
import plotly.io as pio
pio.renderers.default = "notebook"
pd.options.plotting.backend = 'plotly'

# from dsc259r_utils import * # Feel free to uncomment and use this.
```

Step 1: Introduction

## What features predict recipe popularity (number of ratings)?

Popularity is defined as the **number of non-zero ratings** a recipe receives in the interactions table. This focuses on engagement/visibility rather than the average star rating.

```
In [2]: # TODO
```

## Step 2: Data Cleaning and Exploratory Data Analysis

```
In [3]: import pandas as pd
import numpy as np
import ast
import plotly.express as px
import plotly.io as pio

pio.renderers.default = "notebook"
pd.options.plotting.backend = "plotly"

recipes = pd.read_csv("RAW_recipes.csv")
interactions = pd.read_csv("RAW_interactions.csv")
```

```
print("recipes:", recipes.shape)
print("interactions:", interactions.shape)

recipes.head()
```

recipes: (83782, 12)  
interactions: (731927, 5)

Out [3]:

		<b>name</b>	<b>id</b>	<b>minutes</b>	<b>contributor_id</b>	<b>submitted</b>	<b>tags</b>	<b>nutrition</b>
0	1	brownies in the world best ever	333281	40	985201	2008-10-27	['60-minutes-or-less', 'time-to-make', 'course...']	[138.4, 10.0, 50.0, 3.0, 3.0, 19.0, 6.0]
1	1 in canada	chocolate chip cookies	453467	45	1848091	2011-04-11	['60-minutes-or-less', 'time-to-make', 'cuisin...']	[595.1, 46.0, 211.0, 22.0, 13.0, 51.0, 26.0]
2	412	broccoli casserole	306168	40	50969	2008-05-30	['60-minutes-or-less', 'time-to-make', 'course...']	[194.8, 20.0, 6.0, 32.0, 22.0, 36.0, 3.0]
3	millionaire pound cake		286009	120	461724	2008-02-12	['time-to-make', 'course', 'cuisine', 'prepara...']	[878.3, 63.0, 326.0, 13.0, 20.0, 123.0, 39.0]
4	2000	meatloaf	475785	90	2202916	2012-03-06	['time-to-make', 'course', 'main-ingredient', ...]	[267.0, 30.0, 12.0, 12.0, 29.0, 48.0, 2.0]

In [4]:

```
# --- Parse / clean recipes ---
# Parse list-like columns stored as strings
for col in ["tags", "nutrition", "ingredients", "steps"]:
    if col in recipes.columns:
        recipes[col] = recipes[col].apply(ast.literal_eval)
```

```
# Convert dates
recipes["submitted"] = pd.to_datetime(recipes["submitted"], errors="co

# Expand nutrition vector into columns (7 values)
nutrition_cols = ["calories", "total_fat", "sugar", "sodium", "protein",
                   "carbohydrates", "fiber"]
nutrition_df = pd.DataFrame(recipes["nutrition"].to_list(), columns=nutritio

recipes_clean = pd.concat([recipes.drop(columns=["nutrition"]), nutrition_d

# Log cook time (very skewed)
recipes_clean["log_minutes"] = np.log1p(recipes_clean["minutes"])

recipes_clean[["id", "name", "minutes", "n_steps", "n_ingredients", "submitted", "calories", "total_fat", "sugar", "sodium", "protein", "carbohydrates", "fiber", "log_minutes"]]

```

Out [4] :

	<b>id</b>	<b>name</b>	<b>minutes</b>	<b>n_steps</b>	<b>n_ingredients</b>	<b>submitted</b>	<b>calories</b>	<b>total_fat</b>	<b>sugar</b>	<b>sodium</b>	<b>protein</b>	<b>carbohydrates</b>	<b>fiber</b>	<b>log_minutes</b>
0	333281	1 brownies in the world best ever	40	10	9	2008-10-27	138.4							
1	453467	1 in canada chocolate chip cookies	45	12	11	2011-04-11	595.1							
2	306168	412 broccoli casserole	40	6	9	2008-05-30	194.8							
3	286009	millionaire pound cake	120	7	7	2008-02-12	878.3							
4	475785	2000 meatloaf	90	17	13	2012-03-06	267.0							

In [5] : # --- Parse / clean interactions ---

```
interactions["date"] = pd.to_datetime(interactions["date"], errors="co

# Remove placeholder/non-ratings (rating == 0)
interactions_clean = interactions[interactions["rating"] > 0].copy()

# Popularity = number of (non-zero) ratings per recipe
popularity = (
    interactions_clean
        .groupby("recipe_id")
        .size()
        .reset_index(name="num_ratings")
)
```

```
# Merge popularity back to recipe-level table
recipe_level = recipes_clean.merge(popularity, left_on="id", right_on="id")
recipe_level["num_ratings"] = recipe_level["num_ratings"].fillna(0).as_type("float64")

# Log popularity (very skewed)
recipe_level["log_num_ratings"] = np.log1p(recipe_level["num_ratings"])

recipe_level[["id", "name", "num_ratings", "log_num_ratings", "minutes", "log_minutes", "n_rated"]]
```


Out[5]:

	<code>id</code>	<code>name</code>	<code>num_ratings</code>	<code>log_num_ratings</code>	<code>minutes</code>	<code>log_minutes</code>	<code>n_rated</code>
0	333281	1 brownies in the world best ever	1	0.693147	40	3.713572	1
1	453467	1 in canada chocolate chip cookies	1	0.693147	45	3.828641	1
2	306168	412 broccoli casserole	4	1.609438	40	3.713572	412
3	286009	millionaire pound cake	1	0.693147	120	4.795791	1
4	475785	2000 meatloaf	2	1.098612	90	4.510860	2000

In [6]:


```
# Univariate: popularity distribution (log scale)
fig = px.histogram(
    recipe_level,
    x="log_num_ratings",
    nbins=60,
    title="Log Distribution of Recipe Popularity (Number of Ratings)"
)
fig.show()
```

## Log Distribution of Recipe Popularity (Number of Ratings)



```
In [7]: # Univariate: ingredient count distribution
fig = px.histogram(
    recipe_level,
    x="n_ingredients",
    nbins=50,
    title="Distribution of Number of Ingredients"
)
fig.show()
```

## Distribution of Number of Ingredients



```
In [8]: # Univariate: cook time distribution (log scale)
fig = px.histogram(
    recipe_level,
    x="log_minutes",
    nbins=60,
    title="Log Distribution of Cook Time (minutes)"
)
fig.show()
```


## Log Distribution of Cook Time (minutes)



```
In [9]: # Bivariate: ingredients vs popularity
sample = recipe_level.sample(15000, random_state=42)


fig = px.scatter(
    sample,
    x="n_ingredients",
    y="log_num_ratings",
    trendline="ols",
    title="Ingredient Count vs Popularity (log #ratings)"
)
fig.show()
```

## Ingredient Count vs Popularity (log #ratings)




```
In [10]: # Bivariate: cook time vs popularity
fig = px.scatter(
    sample,
    x="log_minutes",
    y="log_num_ratings",
    trendline="ols",
    title="Cook Time vs Popularity (log #ratings)"
)
fig.show()
```

## Cook Time vs Popularity (log #ratings)



```
In [11]: # Bivariate: calories vs popularity
fig = px.scatter(
    sample,
    x="calories",
    y="log_num_ratings",
    trendline="ols",
    title="Calories vs Popularity (log #ratings)"
)
fig.show()
```

## Calories vs Popularity (log #ratings)



```
In [12]: # Interesting aggregates: popularity by ingredient bucket
recipe_level["ingredient_bucket"] = pd.cut(
    recipe_level["n_ingredients"],
    bins=[0, 5, 10, 15, 20, 50, np.inf],
    labels=["1-5", "6-10", "11-15", "16-20", "21-50", "50+"],
    include_lowest=True
)

agg_ing = (
    recipe_level
    .groupby("ingredient_bucket", observed=True)[["num_ratings"]]
    .agg(["mean", "median", "count"])
    .reset_index()
)

fig = px.bar(
    agg_ing,
    x="ingredient_bucket",
    y="mean",
    title="Average Popularity (mean #ratings) by Ingredient Bucket",
```

```
        hover_data=["median", "count"]
    )
fig.show()

agg_ing
```

## Average Popularity (mean #ratings) by Ingredient Bucket



Out[12]:


ingredient_bucket	mean	median	count
0	2.816718	2.0	14164
1	2.637028	1.0	41623
2	2.496317	1.0	22810
3	2.431808	1.0	4502
4	2.704246	1.0	683

In [13]:

```
# Interesting aggregates: popularity by cook time bucket
recipe_level["time_bucket"] = pd.cut(
    recipe_level["minutes"],
    bins=[0, 15, 30, 60, 120, 240, np.inf],
```

```
    labels=["<=15", "15–30", "30–60", "60–120", "120–240", "240+"],  
    include_lowest=True  
)  
  
agg_time = (  
    recipe_level  
    .groupby("time_bucket", observed=True) ["num_ratings"]  
    .agg(["mean", "median", "count"])  
    .reset_index()  
)  
  
fig = px.bar(  
    agg_time,  
    x="time_bucket",  
    y="mean",  
    title="Average Popularity (mean #ratings) by Cook Time Bucket",  
    hover_data=["median", "count"]  
)  
fig.show()  
  
agg_time
```

Average Popularity (mean #ratings) by Cook Time Buc



Out [13]:

	time_bucket	mean	median	count
0	<=15	2.781355	2.0	16680
1	15-30	2.686264	1.0	20632
2	30-60	2.586245	1.0	25416
3	60-120	2.413774	1.0	12328
4	120-240	2.357459	1.0	4297
5	240+	2.699932	1.0	4429

## Step 3: Assessment of Missingness

In [14]: # TODO

## Step 4: Hypothesis Testing

In [15]: # TODO

## Step 5: Framing a Prediction Problem

In [16]: # TODO

## Step 6: Baseline Model

In [17]: # TODO

## Step 7: Final Model

In [18]: # TODO

## Step 8: Fairness Analysis

In [19]: # TODO